

Fixed Lag Particle Filtering for Target Tracking

Pete Bunch

August 2011

A report on the first year's progress



UNIVERSITY OF
CAMBRIDGE

Contents

1	Introduction	5
2	Literature Review	7
2.1	Particle filters	7
2.1.1	Sequential importance sampling	8
2.1.2	Markov chain Monte Carlo	9
2.1.3	Coping with dimensionality	11
2.1.4	Marginalised particle filters	11
2.2	Tracking	12
2.2.1	Probabilistic data association	12
2.2.2	Data association hypothesis methods	14
2.2.3	Full particle filters	15
2.2.4	Probability hypothesis density methods	15
2.2.5	Bringing it all together	15
3	Mathematical foundations	17
3.1	Bayes of our lives - A short introduction to filtering	17
3.2	Keep Kalman carry on - the Kalman filter and its extensions	18
3.2.1	The basic filter	18
3.2.2	The extended filter	19
3.2.3	The Kalman smoother	19
3.3	Tough as old bootstraps - the basic particle filter	19
3.3.1	The basics	20
3.3.2	Auxiliary sampling	21
3.3.3	Degeneracy and resampling	22
3.3.4	Importance distributions	22
3.4	MCMC in da house - Markov chains for particle filtering	23

4	The tracking model	25
4.1	Observation Likelihood	26
4.2	Transition Dynamics	26
4.3	Association Likelihood	26
4.4	Assumptions	27
5	Fixed Lag Estimation for Tracking	29
5.1	The benefits of delaying	29
5.2	A mathematical framework for fixed lag estimation	30
5.2.1	Proposing older histories	31
5.3	Application to tracking	31
5.3.1	Importance distributions	32
6	Fixed Lag Particle Filters	35
6.1	Sequential Importance Sampling and Resampling	35
6.1.1	Coping with dimensionality	35
6.2	Markov Chain Monte Carlo	35
6.3	Marginalised Particle Filters	35
7	Plan Of Future Research	37

Introduction

In a target tracking scenario, the aim is to trace the trajectory of an object over time from a set of discrete observations.

filler here

The additional difficulty in target tracking is the nature of the observation process. Our imperfect sensors may not detect every target present in the scene in every scan. Furthermore, there may be some number of false alarms arising from sensor errors or clutter. Our task thus becomes three-fold: firstly to detect what targets are present in the scene, secondly to work out which observations were generated by each target, and finally to estimate the states of the targets. The basic Kalman and particle filters address only the third of these tasks.

Research in target tracking emerged from military applications such as radar and sonar. However, similar problems emerge in many diverse branches of science: the tracking of people, vehicles or animals in video sequences; of molecules or cells in microscopy data; or of notes in a piece of music. Each can be reduced to a similar underlying model. In this chapter we will formulate such a basic model, keeping it as general as possible rather than focusing on any specific application.

Report Structure

Literature Review

Two spheres of literature will be of interest here: the developments of models and algorithms for tracking, and the studies of the numerical inference schemes which underly our method. We begin with the latter.

2.1 Particle filters

Filtering is the operation of inferring the state of an evolving latent random process given a set of imperfect observations up until the current time. In a Bayesian estimation scheme, we would like to calculate a posterior probability distribution for such a state at each point in time. The difficulty with such a procedure is that the complexity of the state distribution will tend to compound over time. In the particular case where both the state transition and observation processes are linear transformations with Gaussian noise, it is possible to derive an analytic expression for the state distribution with constant complexity, the fabulous Kalman filter (KF) of (Kalman 1960). However, as many problems are often non-linear or non-Gaussian, and as no other such analytic and generally applicable cases have been discovered, (Daum 2005), we must resort to numerical techniques.

The particle filter (PF) approximates a probability distribution using a set of discrete samples or “particles” drawn from it. This will be no trivial task - the distribution cannot be sampled directly, just as it cannot be expressed analytically. Instead, at each time step, the cloud of particles representing the previous distribution is propagated forwards to approximate the next, taking into account the latest observation. The most common PF algorithm in use is known as Sequential Importance Sampling with Resampling (SISR) (or some permutation of this initialism) - in fact the terms SISR and PF are often used interchangeably.

2.1.1 Sequential importance sampling

Importance sampling

The first modern implementation of the SIS algorithm by (Gordon, Salmond & Smith 1993). At the heart of the algorithm is an importance sampling step. Each particle from the old distribution is propagated forward by sampling from a proposal or ‘importance’ distribution. The particles are then assigned an “importance weight” to account for the discrepancy between the proposal distribution and the desired posterior. A significant strength of the algorithm is that we need only be able to calculate the posterior probability up to a normalising constant. This may be easily accounted for by normalising the weights of the discrete sample set.

Resampling

Over time, the variance in these weights is liable to grow, as unlikely particles deviate completely from the correct state and are assigned a negligible weight. The eventual result is a degenerate sample with all the weight on one or a few particles and most contributing nothing to the approximation. The solution to this degeneracy problem is to introduce a resampling step before the importance sampling, in which particles are sampled from the particle approximation. Thus, heavily weighted particles are multiplied while low weight particles are discarded. Many algorithms for resampling exist, a summary of which is given by (Doucet & Johansen 2009). This resampling step introduces a degree of Monte Carlo error, and so (Liu & Chen 1995) introduce a measure of particle diversity to assess the degree of degeneracy. This measure, the effective sample size (ESS), is an estimate of the equivalent number of equal-weighted particles in the approximation, and is based on a calculation of the sample weight variance. Resampling is only used if the ESS falls below a chosen threshold.

Importance distributions

The filter proposed by (Gordon et al. 1993), known as the bootstrap filter, used the state transition density as the importance distribution, as did many other early implementations, (Blake, Isard et al. 1998, Kitagawa 1996). However, this can result in poor performance if the observation process is informative. In other words, if we know that the observation is always very close to the state, then there is no point in proposing particles far away which will have a very low weight. It was suggested in, amongst others, (Liu & Chen 1995) that a better proposal could be constructed by considering the position of the new observation. In particular, the “optimal” (in the sense that it minimises the weight variance) proposal is the conditional distribution of the new state given the state history and the new observation (see (Doucet, Godsill & Andrieu 2000) for proof). This is rarely available analytically or samplable, so Gaussian approximations may be used (Doucet et al. 2000).

Auxiliary particle filtering

In the description above, resampling was introduced as an additional procedure to rejuvenate a degenerate particle distribution. However, it is possible to view it as a more integral part of the sampling procedure: At each step we construct a set of particles by proposing both a new state and a history for each. These histories are selected from the previous particle distribution. The

resampling fulfils the role of a “history proposal”. On a step with resampling we propose histories from the weighted particle distribution from the previous frame. On a step without resampling, we propose histories from the unweighted particle distribution, which can of course be done by simply keeping the same set of particles! In this new paradigm the obvious generalisation is to allow history proposals from the previous set of particles with arbitrary weights. A correction is made in the importance weight calculation to account for the proposal weights. This is the innovation of the auxiliary particle filter introduced by (Pitt & Shephard 1999). The proposal weights may be assigned so as to favour particles which give rise to more likely states.

Resample-move

The objective when selecting the “optimal” proposal distribution or in choosing auxiliary proposal weights is to minimise the variance of the importance weights, and thus reduce the need for resampling. Another important particle filtering technique is the resample-move method of (Gills & Berzuini 2001), which aims instead to replenish the particle diversity after resampling. Rather than beginning the next IS step with multiple copies of the same particle, the current and previous states of each are adapted using one or more Metropolis-Hastings moves. Although this method improves the representation of the state distribution, it requires significant additional computation for each particle.

2.1.2 *Markov chain Monte Carlo*

Metropolis-Hastings

Markov chain Monte Carlo (MCMC) algorithms, like SISR, make use of a set of samples to approximate a probability distribution. Unlike SISR, MCMC is not specifically designed to address problems of sequential inference, although, as we shall see, they may be applied to these problems as an MCMC particle filter.

The basic MCMC method was proposed by (Metropolis, Rosenbluth, Rosenbluth, Teller, Teller et al. 1953) and extended by (Hastings 1970), after which the Metropolis-Hastings (MH) algorithm is named. Rather than producing independent samples from a distribution, MH uses a Markov chain. A reversible, ergodic Markov chain has a stationary distribution. By careful design of the transition properties of the chain, we can make the stationary distribution equal to some desired target distribution. For our inference problems, this target will be the state posterior.

A MH move is made by sampling a value of the state from a proposal distribution. With some probability, the new value is accepted, otherwise the previous value is maintained. The acceptance probability is given by:

$$\alpha = \min \left(1, \frac{P(X_{\text{new}})q(X_{\text{old}}|X_{\text{new}})}{P(X_{\text{old}})q(X_{\text{new}}|X_{\text{old}})} \right) \quad (2.1)$$

where $q(\cdot|\cdot)$ is the proposal distribution, and $P(\cdot)$ the target.

The chain must be initialised at some value, and this is unlikely to be a sample from the target (if we could sample it, we would not need MCMC). Thus, there will be a period over which the sampler converges to the target distribution. This period is known as the “burn in” and the samples should be excluded from the particle approximation.

The choice of proposal does not affect the target distribution, although it will affect the rate at which the algorithm converges. If the proposal distribution has a low variance then successive samples will be close together, and it may take many moves to explore the whole distribution. If the variance is large, then the acceptance probability may be very low and again convergence may be slow.

A complete and excellent discussion of MCMC design considerations, convergence properties, etc. can be found in (Gilks, Richardson & Spiegelhalter 1996), and will not be repeated here.

Block sampling

When MCMC algorithms are applied to multivariate distributions it is not necessary to change every variable in every move. If the state is divided into two blocks, X_i and X_{-i} , and a move is proposed only on X_i , then the acceptance probability is now given by:

$$\alpha = \min \left(1, \frac{P(X_{i,\text{new}}|X_{-i})q(X_{i,\text{old}}|X_{i,\text{new}}, X_{-i})}{P(X_{i,\text{old}}|X_{-i})q(X_{i,\text{new}}|X_{i,\text{old}}, X_{-i})} \right) \quad (2.2)$$

A special case of this scheme occurs when $q(X_{i,\text{new}}|X_{i,\text{old}}, X_{-i}) = P(X_{i,\text{new}}|X_{-i})$. This is the Gibbs sampler devised by (Geman & Geman 1984). With this choice of proposal distribution, the acceptance probability cancels out to 1. Thus, all samples are accepted. The difficulty with this formulation is that the required conditional distribution may not be available, or may not be samplable.

MH moves which alter only a subset of the variables are commonly known as Metropolis-within-Gibbs (MwG) moves.

Reversible jump Markov chain Monte Carlo

MCMC methods are not restricted to distributions of fixed dimension. The Reversible Jump MCMC (RJMCMC) method devised by (Green 1995) allows MH moves to jump between state spaces with different dimensionality. This is especially useful when trying to choose between models while simultaneously estimating parameters, or for estimating the order or number of components in a model, for example in a Gaussian mixture or autoregressive model.

RJMCMC moves may have very low acceptance probabilities. When the sampler jump from one state space to another, there may be no obvious function for generating likely values in the new state. For such cases, (Al-Awadhi, Hurn & Jennison 2004) propose using a system of compound moves, where a model change move is followed by a number of normal MH moves designed to seek out a likely set of state values. A single acceptance probability is then calculated for the set of moves.

Sequential Markov chain Monte Carlo

The interest in MCMC for this project lies in their application to sequential inference problems. Such ‘MCMC particle filters’ were first proposed by (Khan, Balch & Dellaert 2005), although this first implementation suffered from computation issues, and (Golightly & Wilkinson 2006), and have recently been applied to a range of problems including target tracking, (Septier, Pang, Carmi & Godsill 2009), group tracking, e.g. (Carmi, Godsill & Septier 2009), tracking of a correlated stock prices, (Pang, Godsill, Li & Septier 2011) and music transcription, (Bunch & Godsill 2010).

MENTION POLSON'S PRACTICAL FILTER SOMEWHERE - try to understand it first.

The MCMC particle filter runs a Markov chain for each frame which targets the posterior state distribution at that time. The particles are thus unweighted and are generated in series rather than in parallel, with each particle depending on the last. The state space of the Markov chain consists of the entire history of the hidden variables as well as the latest values. Two types of MH move are thus required: History-moves involve proposing a new history from the particles of the previous particle distribution. Current-moves involve a proposal of the latest variables given the history and data. The two types of move are analogous to the resampling and importance sampling steps of the SISR algorithm respectively.

MCMC particle filters have a significant advantage over their SISR counterparts for complex, high dimensional problems. As noted in (Pang, Li & Godsill 2008), we can use MwG moves, changing only a few variables at once, to explore the posterior state distribution. In contrast, with an SISR particle filter, we must propose a new value for every variable at once. In some cases this enables the MCMC particle filter to give a better approximation of the state posterior, as demonstrated in (Pang et al. 2011).

2.1.3 *Coping with dimensionality*

It is a commonly noted problem that the efficacy of particle filters falls as the number of state dimensions increases. This is a problem for both SISR and MCMC filters, and indeed for any Monte Carlo approximation. There is a strong intuition for this phenomenon - it will clearly take more particles to adequately sample a square than it would a line, and more still for a cube. An analysis of this effect, including numerical studies, was conducted in (Daum & Huang 2003).

The severity of the “curse of dimensionality” will depend on the correlation between the state variables, and the ability to exploit such correlations by the design of effective proposal distributions. Methods to improve performance of MCMC algorithms include Hybrid Monte Carlo (HMC) (Duane, Kennedy, Pendleton & Roweth 1987), which introduces “momentum” variables to the state in order to build better proposals and the Riemann manifold methods of (Girolami & Calderhead 2011).

2.1.4 *Marginalised particle filters*

Monte Carlo methods are effective for tackling nonlinear, non-Gaussian problems when analytic methods cannot be found. However, they are computationally expensive and so should be avoided when an analytic solution does exist. In some problems, the state can be partitioned into two parts, one which behaves in a linear-Gaussian way conditional on the other. In such a case a particle filter can be used for inference of the nonlinear part, and a Kalman filter for the linear-Gaussian part. Such a strategy reduces the variance of the state estimates, in accordance with the Rao-Blackwell theorem. See, for example, (Casella & Robert 1996). The only change required is that likelihood calculation for the particle filter will now require the linear-Gaussian part of the state to be marginalised. Such schemes may be used with both SISR and MCMC particle filters and are known as *marginalised* or *Rao-Blackwellised* particle filters.

2.2 Tracking

In a target tracking situation we aim is to trace the trajectory of an object over time from a set of discrete observations. For the simplest case, with a single target under observation, which is detected in every frame, and with no false alarms, this is a simple state-space inference problem. It can be approached with a Kalman filter or a basic particle filter. The phenomenon which makes tracking problems more challenging is association ambiguity. In a given frame, we may not be guaranteed to detect a target, and we may pick up false alarms or clutter measurements. In general there is no way to establish with certainty which observation arose from which target. Finally, the number of targets in the scene may also be unknown. We are thus faced with a three-fold problem:

- Detect how many targets are present
- Estimate which observation arose from each target
- Estimate the state of each target

In this section, we review a number of strategies for tackling such a target tracking problem. Note that although we divide up the algorithms into sections, there is significant overlap between them, which we endeavour to highlight.

2.2.1 Probabilistic data association

The earliest works on target tracking with data association used a combination of Kalman filters and heuristics. (Sea 1971) suggests using only the observation with the minimum Mahalanobis distance from the Kalman prediction, i.e. a maximum likelihood estimate of the association. Such nearest neighbour algorithms can easily be caused to lose track by a single unfortunate false alarm.

A probabilistic approach was introduced by (Bar-Shalom & Tse 1975), in the form of the probabilistic data association filter (PDAF). Rather than select a single observation for each target in each frame, for the PDAF a posterior probability of each possible association is calculated, given the previous target states. A Kalman update is then evaluated for each of the possible associations and the results added together, weighted by the corresponding association posterior. Thus, the final update takes into account multiple (indeed, all) observations. Thus the PDAF cannot be so easily upset by single clutter measurement, although additional error is introduced in situations where a nearest neighbour method would have picked the correct association, because of the influence of false alarms which now contribute to the estimate.

PICTURE - demonstrating PDAF

When more than one target is present in the scene, a separate PDAF can be run for each one. However, this leads to errors. When two tracks pass close together, they may both have a high association probability with the same observation(s). The result is that the estimates converge onto the same path, and one of the tracks is lost.

PICTURE - demonstrating track merging

The shortcomings of the PDAF are addressed by extending the filter to consider the joint association posteriors. This is the joint PDAF (JPDAF) of (Fortmann, Bar-Shalom & Scheffe 1983), excellently reviewed in (Bar-Shalom, Daum & Huang 2009). Rather than calculating association probabilities for each target in isolation, the joint association hypotheses are assessed. We can now include as prior information the fact that no observation can be associated with two targets. Thus the problem of tracks following the same observations is reduced. The price

of this improved performance is the need to calculate an association probability for every joint hypothesis. The number of these is of combinatorial complexity in the number of targets and observations in a frame. Gating is used to address this issue: a target-observation association is only considered possible if the Mahalanobis distance between the two is below some threshold. The choice of threshold governs the probability of excluding the correct hypothesis, and thus trades off performance and complexity (Sea 1971).

Both the PDAF and the JPDAF suffer from a problem of bias and track coalescence when targets pass close to each other. This was observed and quantified by (Fitzgerald 1985). The measurements generated by the second target ‘pull’ the estimate of the first away from its correct position. When two targets are travelling along parallel paths this can lead to *track coalescence*. The estimates of both target states move together, midway between the two correct locations. Solutions to this problem often resort to reintroducing nearest-neighbour methods, such as the nearest-neighbour-JPDAF (Fitzgerald 1986) or the set-JPDAF (Svensson, Svensson & Willett 2009).

So far we have outlined how the JPDAF approaches the problems of data association and state estimation. It remains to consider how target detection may be incorporated. One of the first methods proposed was that of (Bar-Shalom, Chang & Blom 1989) which used the interacting multiple model (IMM) algorithm in parallel with the JPDAF. When a the possibility of a new target existed, two filters were run in parallel, one assuming the existence of the new target, and one assuming its non-existence. Decision logic was included to choose between the two possibilities. A less computationally demanding approach was formulated by (Musicki, Evans & Stankovic 1994, Musicki & Evans 2004), in which each target was labeled with a probability of existence, allowing potential new targets to be assessed in parallel with tracking.

The JPDAF requires the dynamics of the targets and the observation process to be linear and Gaussian, or that an appropriate EKF-like linear approximation can be made. However, even given this requirement, the estimation process still involves a sub-optimal step. The final posterior state distribution is constructed by adding a weighted sum of distributions, each made with a different association hypothesis. The individual components are Gaussians, and so the posterior should be a sum of Gaussian. However, in the JPDAF, this is collapsed into a single Gaussian with a matched mean and variance. This collapsing step may throw away important information if there is more than one significant component, and is responsible for the track coalescence effect noted before. It is possible to construct a filter which maintains the complete or a partially complete Gaussian sum posterior, (Singer, Sea & Housewright 1974, Salmond 1990), at the expense of increased complexity. In fact, this produces an algorithm broadly equivalent to the multi-hypothesis tracker (Blackman 2004)!

There is an alternative to using a Kalman filter for the JPDAF, with its associated approximations. We can use a particle filter instead for the state estimation. Such an algorithm is developed in (Vermaak, Godsill & Perez 2005). A set of particles representing targets states are propagated forwards using IS and used to calculate the association priors using a Monte Carlo estimate. These association priors are then used in the particle weight calculations. The particles approximate the target posterior state distribution. Using a particle approximation allows us to trade off accuracy against computation by varying the number of particles, which may be an improvement on Gaussian approximations for nonlinear models.

A final note about the JPDAF. The algorithm assumes that the only interaction between targets is through the associations - no two targets can be associated with the same observation. Thus, once the association probabilities have been calculated, each target may be updated

independently using a Kalman filter to give a marginal state distribution. Interactions between targets are difficult to accommodate because we do not work in the joint state space of all targets.

2.2.2 Data association hypothesis methods

With a JPDAF, the probability of each feasible combination of associations is calculated, and used to weight a sum of the state estimates made with the respective combinations. The Multi-Hypothesis Tracker (MHT) of (Reid 1979) instead maintains each of these estimates separately, with an associated probabilistic score. The MHT requires the target dynamics to be exactly or approximately linear-Gaussian, so that KFs can be used for state updates. As we shall see, a PF version results in Monte Carlo Data Association.

For each frame of data, each possible combination of associations between objects and observations is formed. State estimates are calculated with KF updates and the hypothesis is scored with a posterior probability. Detection of new tracks is readily incorporated into MHT. New tracks can be added into hypotheses wherever an un-associated observation exists. MHT has a potentially enormous computational complexity which grows combinatorially as time proceeds. To render it practical, observations are gated, disallowing associations between targets and distant measurements, and ‘pruning’ is used to eliminate hypotheses with low probabilities. (Blackman 2004) contains an excellent introduction to MHT and details of its implementation.

MHT requires significantly more computational complexity than the JPDAF. In particular, high levels of clutter can lead to prohibitively large numbers of hypotheses. Furthermore, it is reliant on the use of Gaussian approximations so that KF methods can be used to obtain state estimates.

A modification of MHT is derived by relaxing the constraint that any observation can only be associated with one target. This allows the tracking of each target to be conducted independently, with a great saving in computation (Complete lists of hypotheses are no longer required). The EM algorithm may be used to select the maximum probability association hypothesis over a window of frames. This is known as Probabilistic MHT (PMHT) and was first proposed by (Streit & Luginbuhl 1994). (Willett, Ruan & Streit 2002) show that performance is similar to that of the PDAF (The PMHT assumption is equivalent to running an independent PDAF on each target instead of a JPDAF on the whole lot).

MHT suffers from high computational loads due to the large number of possible association hypotheses. A particle method was introduced to address this problem in (Oh, Russell & Sastry 2004), called MCMC data association (MCMCDA). This method still uses Gaussian approximations so that KFs may be used to estimate the state distributions. However, instead of enumerating a posterior probability for every possible hypothesis, a Markov chain is constructed to target the posterior association hypothesis probability. This is a batch method, allowing changes to the associations in previous frames within some fixed length window. MH moves allow transitions between valid hypotheses, for example by initiating new tracks, extending or shortening tracks, or swapping observations between nearby targets. By choosing sensible proposals, low probability hypotheses are never even considered. In (Oh et al. 2004, Oh, Russell & Sastry 2009), the authors report significant improvements over MHT.

2.2.3 Full particle filters

In the MCMCDA method, a particle distribution was developed over the associations. In the MC-JPDAF, a particle filter was used for the target posterior state distributions. The next step is to maintain a particle distribution over both states and associations of all targets. Such a method using SISR was proposed by (Vermaak et al. 2005). The targetted distribution of the particle filter is the joint posterior of the target states and associations. Both these components must be proposed for each target, and the probability of each used in the weight updates.

Full multi-target particle filters suffer from complexity issues. The dimensionality of the state space scales with the number of targets. Thus the variance of importance weights, or acceptance probabilities, increases rapidly. An SISR particle filter with more than a couple of targets may repeatedly have all the weight on a single particle, even with resampling every step. An intuitive interpretation of the effect is that a particle may be given a low weight because its estimate of one target state is poor, even if the others are all good. In (Vermaak et al. 2005), two strategies are proposed to cope with the dimensionality problems. In the Sequential Sampling Particle Filter (SSPF), targets are sampled sequentially, with optional resampling steps between each. This can be used to maintain a higher level of particle diversity, but there is still a worse than linear scaling in complexity to maintain a constant level of accuracy. For the Independent Partition Particle Filter, targets are assumed to be completely independent, by allowing multiple targets to associate with the same observation, as with PMHT. This allows an independent particle filter to be run for each target from which the joint distribution may be reconstructed.

Interestingly, the a marginalised or Rao-Blackwellised version of the full multi-target particle filter (i.e. approximating the target states with Kalman filters) is equivalent to the MCMCDA approach. Such an algorithm using SISR is presented in (Särkkä, Vehtari & Lampinen 2007).

2.2.4 Probability hypothesis density methods

All the methods outlined so far assume that each target is identified by some unique label. There exists another family of tracking algorithms based on the assumption that the multi-target state is an unordered set, i.e. that it does not matter which target is which, only where they occur. The finite set statistics (FISST) required to handle states which are random finite sets (RFS) was presented by (Mahler 1994). Practical algorithms based on RFS methods are generally restricted to estimating the first moment of the multitarget probability distribution, known as the “probability hypothesis density” (PHD), (Mahler 2003), which can be approximated using Gaussian mixtures, (Vo & Ma 2006), or particle filters, (Vo, Singh & Doucet 2005, Whiteley, Singh & Godsill 2010). A more detailed introduction to PHD methods can be found in (Mahler 2004, Wood 2010).

2.2.5 Bringing it all together

The methods we have considered can broadly be divided into those which attempt to explicitly estimate the associations between targets and observations, and those that marginalise this information and estimate only a combined state estimate. We can also divide the methods into those which use Gaussian approximations and Kalman filters, and those which use particle approximations for the target states.

	Estimate Associations	Marginalise Associations
Gaussian State Approximation	MHT MCMCDA RBPF	JPDAF GM-PHD
Particle State Approximation	Full PF	MC-JPDAF SMC-PHD

Table 2.1: Loose grouping of target tracking algorithms by the treatment of associations and state approximations

Mathematical foundations

3.1 Bayes of our lives - A short introduction to filtering

For the literature review, mathematical detail was kept to a minimum for clarity of presentation. In this chapter we revisit some basics of inference and particle filtering in order to establish notation and mathematical foundations for the later chapters.

Many tasks in signal processing, science in general, and indeed life, require us to make some estimate of an unknown quantity from indirect, incomplete, or inaccurate observations. By constructing a model to explain how these observations depend on the underlying state, we can infer something about that state. We will express this observation model in terms of a likelihood function:

$$P(Y|X) \tag{3.1}$$

where X is the state and Y the observations. This is not the whole story - in many cases we are not estimating our unknown state “from scratch”. Previous experience, prejudice, and prior knowledge can also contribute to our estimates. The likelihood and prior terms can be combined through our friend, Bayes rule (Bayes & Price 1763, Laplace 1774), to calculate the posterior probability of the state, i.e. the probability of the state given the observations:

$$P(X|Y) = \frac{P(Y|X)P(X)}{P(Y)} \tag{3.2}$$

This is the basis of the process of inference. Mathematically, we can assign a probability distribution to the state space of X . By applying Bayes rule, we are updating our belief about the values of X using the information in Y .

Often the quantity in which we are interested, X , is changing over time, and we would like to estimate its value at each point in time given only the observations received so far. In this case we will generally assume that the state is Markovian, i.e. that $X_t|X_{t-1}$ is independent of $X_{1:t-2}$, giving us a hidden Markov state-space model. This is the traditional filtering problem. By constructing a model for the evolution of the unknown state, we can now derive our prior information from our estimate at the previous time step. In discrete time, we now write:

$$P(X_t|Y_{1:t}) = \frac{\int P(Y_t|X_t)P(X_t|X_{t-1})P(X_{t-1}|Y_{1:t-1})dX_{t-1}}{P(Y_t|Y_{1:t-1})} \quad (3.3)$$

where the subscript indicates the time and ranges are notated by : in the MATLAB style. $P(X_t|Y_{1:t})$ is called the filtering distribution. Equation 3.3 describes the ideal Bayesian filter.

So far, we have expressed the problem entirely in terms of distributions. The same models may be expressed in terms of difference equations of random variables. In the most general form:

$$X_t = f_t(X_{t-1}, V_t) \quad (3.4)$$

$$Y_t = g_t(X_t, W_t) \quad (3.5)$$

where f_t and g_t are known deterministic functions and V_t and W_t and random variables, known as the process and observation noise respectively.

3.2 Keep Kalman carry on - the Kalman filter and its extensions

3.2.1 The basic filter

In a few simple cases the filtering set-up permits the derivation of closed form posterior distributions at each time instant. Most notably, the Kalman filter (KF) (Kalman 1960) is an analytic filter for models with continuous state and observation variables, in which both transition and observation models are linear transformations with Gaussian innovations.

$$x_t = Ax_t + v_t \quad (3.6)$$

$$y_t = Cx_t + w_t \quad (3.7)$$

where v_t and w_t are now Gaussian random variables with zero mean and covariance matrices Q and R . We use lower case variables here to emphasise that these are “nice”, continuous vectors. (As we shall see, our state variable will later be sets or lists).

Kalman’s solution for the linear-Gaussian case is given by:

$$P(x_t|y_{1:t}) = \mathcal{N}(x_t|\mu_t, \Sigma_t) \quad (3.8)$$

$$P(x_t|y_{1:t-1}) = \mathcal{N}(x_t|\hat{\mu}_t, \hat{\Sigma}_t) \quad (3.9)$$

where μ_t , Σ_t , etc. are given by the following recursions.

Time Updates:

$$\hat{\mu}_t = A\mu_{t-1} \quad (3.10)$$

$$\hat{\Sigma}_t = A\Sigma_{t-1}A^T + Q \quad (3.11)$$

Measurement Updates:

$$z_t = y_t - C\hat{\mu}_t \quad (3.12)$$

$$S_t = C\hat{\Sigma}_tC^T + R \quad (3.13)$$

$$K_t = \hat{\Sigma}_tC^TS_t^{-1} \quad (3.14)$$

$$\mu_t = \hat{\mu}_t + K_tz_t \quad (3.15)$$

$$\Sigma_t = (I - K_t C) \hat{\Sigma}_t \quad (3.16)$$

The KF is delightful because it not only provides us with a closed-form analytic solution, but the complexity of that solution does not increase as we receive additional measurements. This is a consequence of the fact that the Gaussian distribution is its own conjugate prior. Unfortunately, no other such convenient cases have been discovered (Daum 2005). Analytic solutions to non-linear, non-Gaussian filtering problems generally require unacceptable conditions, such as zero process noise $Q = 0$ (Daum 2005).

3.2.2 The extended filter

Given the loveliness of the KF, the instinct when faced by an intractable non-linear filtering problem is to linearise it. This produces the Extended Kalman Filter (EKF), and is achieved by replacing the A and C matrices in equations 3.10 through 3.16 above with Jacobians:

$$A_t = \left. \frac{\partial f}{\partial x_t} \right|_{\mu_{t-1}} \quad (3.17)$$

$$C_t = \left. \frac{\partial g}{\partial x_t} \right|_{\hat{\mu}_t} \quad (3.18)$$

3.2.3 The Kalman smoother

The KF gives us an optimum estimate of $P(x_t|y_{1:t})$. However, once more data has arrived, we can improve this estimate. For a given set of data, $y_{1:T}$, we can estimate the optimum estimates for all previous state distributions, $P(x_{1:T}|y_{1:T})$ using a Rauch-Tung-Striebel (RTS) smoother, (Rauch, Tung & Striebel 1965). This begins with a normal KF, followed by a backward filtering pass which propagates information to earlier time instances. This backward pass is implement by the following recursions:

$$\tilde{\mu}_t = \mu_t + \Sigma_t A^T \hat{\Sigma}_{t+1}^{-1} (\tilde{\mu}_{t+1} - \hat{\mu}_{t+1}) \quad (3.19)$$

$$\tilde{\Sigma}_t = \Sigma_t + [\Sigma_t A^T \hat{\Sigma}_{t+1}^{-1}] (\tilde{\Sigma}_{t+1} - \hat{\Sigma}_{t+1}) [\Sigma_t A^T \hat{\Sigma}_{t+1}^{-1}]^T \quad (3.20)$$

giving us

$$P(x_t|Y_{1:T}) = \mathcal{N}(x_t|\tilde{\mu}_t, \tilde{\Sigma}_t) \quad (3.21)$$

For a full derivation, see (Rauch et al. 1965). There exist other ways to implement Kalman smoothing in a fixed-interval sense, such as the forward-backward smoother, and in a fixed-lag sense, but they will not be used in this work.

3.3 Tough as old bootstraps - the basic particle filter

In general we will not be so lucky as to have a problem with linear-Gaussian dynamics. In this case, a particle filter (PF) may be the best alternative. The PF is based on the idea that if we cannot represent a probability distribution analytically then we may approximate it as a collection of discrete samples or “particles” drawn from the distribution. Often, it will not even

be possible to sample from the desired distribution. We therefore use importance sampling (IS) or Markov chain Monte Carlo methods to generate these samples. The principle of a particle filter is that we recursively generate such a set of particles to approximate $P(X_{1:t}|Y_{1:t})$ given a previous set representing $P(X_{1:t-1}|Y_{1:t-1})$.

3.3.1 The basics

With a PF, we approximate a probability distribution with a set of (weighted) samples drawn from that distribution.

$$P(X) \approx \frac{1}{N} \sum_m W^{(m)} \delta_{X^{(m)}}(X) \quad (3.22)$$

where $\delta(x)$ represents a unit probability point mass at a point x , and $\sum_m W^{(m)} = 1$. Such a method allows us to represent any probability distribution of arbitrary complexity, including multidimensional, multimodal, mixed distributions. As the number of particles increases, the accuracy of the approximation improves at the expense of computational complexity. Thus we have the required tool for estimation in non-linear, non-Gaussian scenarios.

We still face the problem of how to generate these samples. The conventional method for this is Sequential Importance Sampling with Resampling (SISR). For a more complete and traditional introduction to this method, see (Cappé, Godsill & Moulines 2007) or (Doucet & Johansen 2009). Here we follow an outline similar to that used for the derivation of the auxiliary particle filter of (Pitt & Shephard 1999).

Suppose we have a particle approximation to the joint posterior distribution from the previous frame, $\hat{P}(X_{1:t-1}|Y_{1:t-1})$. Each particle represents a path through time, $X_{1:t-1}^{(m)}$, and has an associated weight, $W_{t-1}^{(m)}$. Let us suppose also that the unweighted particles may be considered samples from another distribution:

$$\mu(X_{1:t-1}|Y_{1:t-1}) \approx \frac{1}{N} \sum_m W^{(m)} \delta_{X_{1:t}^{(m)}}(X_{1:t}) \quad (3.23)$$

Thus for a given particle,

$$\hat{P}(X_{1:t-1}^{(m)}|Y_{1:t-1}) = W_t^{(m)} \mu(X_{1:t-1}^{(m)}|Y_{1:t-1}) \quad (3.24)$$

We would like to generate a new particle set which includes the current time instance. We propose a new set of extended tracks from a factored proposal distribution $X_{1:t} \sim q(X_{1:t}|Y_{1:t}) = q(X_t|Y_t, X_{1:t-1})q(X_{1:t-1}|Y_{1:t})$. The two factors are the proposal probabilities for the new state value, X_t and the history, $X_{1:t-1}$, respectively. Particles are then weighted to take account of the difference between the targeted posterior distribution and the importance distribution:

$$W_t^{(m)} = \frac{P(X_{1:t}^{(m)}|Y_{1:t})}{q(X_{1:t}^{(m)}|Y_{1:t})} \quad (3.25)$$

The simplest choice for the history proposal is $q(X_{1:t-1}|Y_{1:t}) = \mu(X_{1:t-1}|Y_{1:t-1})$, which can be implemented by simply keeping the same set of paths as the previous particle set. This is equivalent to an ordinary IS step with no resampling, and importance weights are given by:

$$W_t^{(m)} = \frac{P(X_{1:t}^{(m)}|Y_{1:t})}{q(X_{1:t}^{(m)}|Y_{1:t})} \approx \frac{W_{t-1}^{(m)} P(X_{1:t}^{(m)}|Y_{1:t})}{P(X_{1:t-1}^{(m)}|Y_{1:t-1}) q(X_t^{(m)}|X_{t-1}^{(m)}, Y_t)} \propto \frac{W_{t-1}^{(m)} P(Y_t|X_t^{(m)}) P(X_t^{(m)}|X_{t-1}^{(m)})}{q(X_t^{(m)}|X_{t-1}^{(m)}, Y_t)} \quad (3.26)$$

Alternatively, we could use $q(X_{1:t-1}|Y_{1:t}) = \hat{P}(X_{1:t-1}|Y_{1:t-1})$ as the history proposal, i.e. sample from the weighted particle distribution which approximates the previous posterior. This is equivalent to an IS step preceeded by resampling. Importance weights are now given by:

$$W_t^{(m)} = \frac{P(X_{1:t}^{(m)}|Y_{1:t})}{q(X_{1:t}^{(m)}|Y_{1:t})} \approx \frac{P(X_{1:t}^{(m)}|Y_{1:t})}{P(X_{1:t-1}^{(m)}|Y_{1:t-1}) q(X_t^{(m)}|X_{t-1}^{(m)}, Y_t)} \propto \frac{P(Y_t|X_t^{(m)}) P(X_t^{(m)}|X_{t-1}^{(m)})}{q(X_t^{(m)}|X_{t-1}^{(m)}, Y_t)} \quad (3.27)$$

3.3.2 Auxiliary sampling

We can generalise the form of our history proposal distribution by weighting the particles from the previous posterior distribution with any arbitrary set of weights.

$$q(X_{1:t-1}|Y_{1:t}) = \frac{1}{N} \sum_m V_t^{(m)} \delta_X(x_{1:t}^{(m)}) \quad (3.28)$$

Now we have

$$\hat{P}(X_{1:t-1}|Y_{1:t-1}) = \frac{W_{t-1}^{(m)}}{V_t^{(m)}} q(X_{1:t-1}|Y_{1:t}) \quad (3.29)$$

giving a general form for the importance weights

$$\begin{aligned} W_t^{(m)} &= \frac{P(X_{1:t}^{(m)}|Y_{1:t})}{q(X_{1:t}^{(m)}|Y_{1:t})} \approx \frac{W_{t-1}^{(m)}}{V_t^{(m)}} \times \frac{P(X_{1:t}^{(m)}|Y_{1:t})}{P(X_{1:t-1}^{(m)}|Y_{1:t-1}) q(X_t^{(m)}|X_{t-1}^{(m)}, Y_t)} \\ &\propto \frac{W_{t-1}^{(m)}}{V_t^{(m)}} \times \frac{P(Y_t|X_t^{(m)}) P(X_t^{(m)}|X_{t-1}^{(m)})}{q(X_t^{(m)}|X_{t-1}^{(m)}, Y_t)} \end{aligned} \quad (3.30)$$

A suitable choice for the auxiliary proposal weights is the predictive likelihood of the next measurement, i.e.

$$V_t^{(m)} = P(Y_t|X_{t-1}^{(m)}) \quad (3.31)$$

This favours the selection of particles which better explain the new observation. This quantity is often not easily calculable, so approximations may be used, for example using a Gaussian approximation.

3.3.3 Degeneracy and resampling

The simplest history proposal is given by $q(X_{1:t-1}|Y_{1:t}) = \mu(X_{1:t-1}|Y_{1:t-1})$. This may be “sampled” by simply keeping the N particles from the previous processing step, which is thus equivalent to importance sampling with no resampling. Multiple steps of this nature will lead to a fall in particle diversity, as many of the weights tend towards 0. This occurs because of the recursive nature of the importance weight calculation in equation ???. The solution to this degeneracy is to use the other form of history proposal, $q(X_{1:t-1}|Y_{1:t}) = \hat{P}(X_{1:t-1}|Y_{1:t-1})$, equivalent to importance sampling with resampling. Such a proposal will require additional computation time, but the weight updates are no longer recursive, see equation ???, so diversity is improved.

The resampling/history-sampling may be conducted in a variety of ways. The simplest is multinomial sampling, in which each particle is sampled independently with replacement. In residual sampling, particles are sampled from the particle distribution without replacement, i.e. after a particle is selected, the probability of selecting it again is reduced. Finally, for systematic resampling, a finite real line is divided into sections corresponding to the weight of each particle. Particles are sampled by choosing regular points along this line, with some constant, random offset. The final method produces the least variance in the number of child particles chosen from each parent given the weights, and thus introduces the least Monte Carlo error.

Degeneracy is measured using the Effective Sample Size of (Liu & Chen 1995), which is given by:

$$ESS = \frac{1}{N} \sum_m W^{(m)2} \quad (3.32)$$

where $W^{(m)}$ are the particle weights, and there are N particles. If ESS falls below some chosen threshold, resampling this required to replenish particle diversity.

3.3.4 Importance distributions

It remains to choose the the importance distribution for the current state, $q(X_t|X_{t-1}, Y_t)$. In the original “bootstrap filter” of (Gordon et al. 1993), this was set equal to the transition density, $P(X_t|X_{t-1})$, which leads to cancellation in the expression for importance weights. This is simple but not necessarily optimal. For example if the process noise is high, the samples of X_t will be widely spread. If, however, the observation noise is comparatively low, many or most of the samples will be far from the observation and will have a low weight, giving us poor particle diversity.

An improved choice of proposal distribution is given by:

$$q(X_t|X_{t-1}, Y_t) = P(X_t|X_{t-1}, Y_t) = \frac{P(Y_t|X_t)}{X_t|X_{t-1}} P(Y_t|X_{t-1}) \quad (3.33)$$

This is often known as the “optimal importance distribution”. The importance weights for a step without resampling are now given by:

$$W_t^{(m)} \propto \frac{W_{t-1}^{(m)} P(Y_t|X_t^{(m)}) P(X_t^{(m)}|X_{t-1}^{(m)})}{q(X_t^{(m)}|X_{t-1}^{(m)}, Y_t)} = 1 \quad (3.34)$$

The sense in which this is an optimal proposal can thus be seen: all weights are equal, and resampling will never be required. In general, the optimal importance distribution will not be samplable, but proposals using Gaussian approximations of this form may be used.

3.4 MCMC in da house - Markov chains for particle filtering

The previous sections have focussed on the traditional, SISR particle filter. Here we examine the MCMC version. The mathematical extensions required are minimal.

MCMC methods allow us to generate an unweighted particle approximation to a target probability distribution. Particles are generated sequentially, each dependent on the last, giving us a Markov chain whose stationary distribution is equal to that which we wish to approximate. For a basic MH step, a proposal distribution is sampled to acquire a new candidate state. With some probability this candidate is accepted as the new state. Otherwise, it is rejected and the old state is kept as the new state. The acceptance probability required to make $P(X)$ the stationary distribution of the Markov chain is given by:

$$\alpha = \min \left(1, \frac{P(X_{\text{new}})q(X_{\text{old}}|X_{\text{new}})}{P(X_{\text{old}})q(X_{\text{new}}|X_{\text{old}})} \right) \quad (3.35)$$

This will be appear familiar as the ratio of the importance weights the old and candidate states would receive in a SISR scheme. In a MCMC particle filtering scheme we target the usual joint posterior distribution:

$$P(X_{1:t}|Y_{1:t}) = \frac{P(Y_t|X_t)P(X_t|X_{t-1})P(X_{1:t-1}|Y_{1:t-1})}{P(Y_t|Y_{1:t-1})} \quad (3.36)$$

As for the SISR particle filter, the proposal distribution factorises as $X_{1:t} \sim q(X_{1:t}|Y_{1:t}) = q(X_t|Y_t, X_{1:t-1})q(X_{1:t-1}|Y_{1:t})$. We can use a biased history proposal in the flavour of the auxiliary particle filter, as described by equation 3.28.

We now adopt the following notation: X for the new state of the Markov chain, and Z for the old state, with $V_{X,t}$ and $V_{Z,t}$ respectively for the auxiliary proposal weights. Using these expansions, the acceptance probability is given by:

$$\begin{aligned} \alpha &= \min \left(1, \frac{P(X_{1:t}|Y_{1:t})q(Z_{1:t}|Y_{1:t})}{P(Z_{1:t}|Y_{1:t})q(X_{1:t}|Y_{1:t})} \right) \\ &= \min \left(1, \frac{P(Y_t|X_t)P(X_t|X_{t-1})q(Z_t|Y_t, Z_{1:t-1})V_{Z,t}}{P(Y_t|Z_t)P(Z_t|Z_{t-1})q(X_t|Y_t, X_{1:t-1})V_{X,t}} \right) \end{aligned} \quad (3.37)$$

The advantage of using a MCMC scheme is that we can vary only some subset of the state variables in each move if we choose. For example, we can have one sort of MH move which keeps the history the same, $X_{1:t} = Z_{1:t}$. This leads to a simplification of the equation for α and an improved probability of acceptance.

The tracking model

In this chapter we introduce the mathematical model for the object tracking problem which will be used throughout the rest of this report.

We first assume that our targets exist at singular points in space, and that targets move independently of each other. At first, we will also assume that we know how many targets are in the scene. Each target will have be characterised by a state vector $x_{j,t}$. This state is composed of continuous-valued coordinates and evolves according to a hidden Markov model (HMM).

$$x_{j,t} = f(x_{j,t-1}, v_{j,t}) \quad (4.1)$$

where $v_{j,t}$ is a random vector.

The targets are observed by a sensor which detects each with a probability P_D . If detected, the sensor returns a point observation at a location $y_t^{(i)}$ given by:

$$y_t^{(i)} = g(x_{j,t}, w_{j,t}) \quad (4.2)$$

where $w_{j,t}$ is another random vector, independent of $v_{j,t}$.

In addition to the target-originating observations, the sensor also detects a number of false alarms. We assume that these are generated by a Poisson process, with a uniform intensity over the observation area. The expected number of clutter observations in a frame is denoted μ_C .

We denote the set of targets present as $X_t = \{x_{1,t}, x_{2,t}, \dots, x_{K_t,t}\}$, and the set of observations as $Y_t = \{y_t^{(1)}, y_t^{(2)}, \dots, y_t^{(M_t)}\}$.

We introduce an association variable for each target, $\lambda_{j,t}$, which indicates which of the observations in that frame was generated by this target. If the target is not detected, then $\lambda_{j,t}$ is set to 0. We denote the set $\Lambda_t = \{\lambda_{1,t}, \lambda_{2,t}, \dots, \lambda_{K_t,t}\}$. As each observation is generated by one target or clutter, no two elements of Λ_t may take the same value, unless 0. This is the constraint that the PMHT and PDAF relax.

If we wished to calculate the posterior distribution over the target states, X_t , we would need to marginalise the association variables in the calculation of the likelihood term:

$$P(Y_t|X_t) = \sum_{\Lambda_t} P(\Lambda_t) \prod_{j=1}^{K_t} P(y_t^{(\lambda_{j,t})}|x_{j,t}) \quad (4.3)$$

where the summation is over all feasible values of Λ_t . The number of terms in this summation combinatorially in the number of targets and the number of observations, and may be prohibitively large if there are many targets or observations in the scene. If it is necessary to calculate the likelihood very often, as in a particle filter, it may be preferable to estimate Λ_t instead of marginalising it. The target distribution is now:

$$P(X_{1:t}, \Lambda_{1:t}|Y_{1:t}) = \frac{P(Y_t|X_t, \Lambda_t)P(X_t|X_{t-1})P(\Lambda_t)P(X_{1:t-1}, \Lambda_{1:t-1}|Y_{1:t-1})}{P(Y_t|Y_{1:t-1})} \quad (4.4)$$

We now consider the terms of this equation one at a time.

4.1 Observation Likelihood

The likelihood given the associations is given by:

$$P(Y_t|X_t, \Lambda_t) = V^{-(M_t - M_{T,t})} \prod_{\lambda_{j,t} \neq 0} P(y_t^{(\lambda_{j,t})}|x_{j,t}) \quad (4.5)$$

where V is the volume of the observation region, and thus V^{-1} is the likelihood of a clutter observation, and where $M_{T,t} = |\{\lambda_{j,t} \neq 0\}|$. This can be conveniently factorised as

$$P(Y_t|X_t, \Lambda_t) = V^{-(M_t - K_t)} \prod_{j=1}^{K_t} \begin{cases} P(y_t^{(\lambda_{j,t})}|x_{j,t}) & \lambda_{j,t} \neq 0 \\ V^{-1} & \lambda_{j,t} = 0 \end{cases} \quad (4.6)$$

This expression contains a factor for each target, and a coefficient independent of the state. This coefficient will either cancel out (MCMC) or be discarded in the weight normalisation (SISR).

4.2 Transition Dynamics

Similarly, because of our independence assumption on the targets, the transition dynamic term will factorise:

$$P(X_t|X_{t-1}) = \prod_{j=1}^{K_t} P(x_{j,t}|x_{j-1}) \quad (4.7)$$

4.3 Association Likelihood

The association variable, λ_t , tells us how many targets have been detected (and how many have not), and how many clutter observations are present. To construct a prior distribution on this, consider the $\lambda_{j,t}$ s in turn. Each one can be either 0 with probability $(1 - P_D)$, or an observation index with probability P_D . If it is the latter, there are $M_t - M_{\text{taken}}$ indexes from which to

choose, where M_{taken} is the number of observations already assigned to a previous target. Once we have considered all the targets, $M_{T,t}$ observations will have been assigned to targets, and all those remaining must be generated by the clutter process. The number of these is Poisson distributed. The association prior for feasible is Λ_t thus given by:

$$P(\Lambda_t) = \underbrace{\frac{P_D^{M_t}}{M_t(M_t-1)\dots(M_t-M_{T,t})}}_{\text{detected targets}} \underbrace{(1-P_D)^{(K_t-M_t)}}_{\text{undetected targets}} \underbrace{\frac{\exp(-\mu_C)\mu_C^{(M_t-M_{T,t})}}{(M_t-M_{T,t})!}}_{\text{clutter}} = \frac{\exp(-\mu_C)\mu_C^{(M_t-M_{T,t})} P_D^{M_t} (1-P_D)^{(1-M_t)}}{M_t!} \quad (4.8)$$

while for invalid Λ_t ($\lambda_{i,t} = \lambda_{j,t}$ for $i \neq j$), the probability is zero. Again, we can factorise this expression over the targets.

$$P(\lambda_t) = \frac{\exp(-\mu_C)\mu_C^{(M_t-K_t)}}{M_t!} \prod_{j=1}^{K_t} \begin{cases} P_D & \lambda_{j,t} \neq 0, \lambda_{j,t} \notin \{\lambda_{1:j-1,t}\} \\ 0 & \lambda_{j,t} \in \{\lambda_{1:j-1,t}\} \\ (1-P_D)\mu_C & \lambda_{j,t} = 0 \end{cases} \quad (4.9)$$

Just as for the observation likelihood, this expression has a factor for each target and a state-independent coefficient which can be omitted for either a SIS or MCMC particle filter.

4.4 Assumptions

Here we summarise the assumptions made in the model formulations:

- Targets and observations occur at singular points.
- The number of targets is known.
- Target states evolve independently of each other according to a HMM.
- Targets are detected in each frame with an independent probability P_D . A detected target generates one observation.
- A single target is associated with each observation.
- Clutter observations are generated by a homogeneous spatial Poisson process.

Fixed Lag Estimation for Tracking

5.1 The benefits of delaying

A basic filter, the Kalman filter included, allows us to estimate the latest value of a hidden state given all previous observations. However, if we wait until the next observation arrives, we can often make a much better estimate - we now know where the state is going as well as where it came from. This is the idea exploited by the Kalman smoother. In fact, as we have seen, the regular particle filter gives us an estimate of the entire state history, $X_{1:t}$. However, due to the resampling process, the particle diversity of the state distribution for previous frames falls as we move further back in time.

PICTURE

As well as giving poor estimates for state distributions, a further disadvantage of this falling diversity in a tracking context is a tendency to lose tracks. Suppose that in some frame the majority of particles associate with a particular likely-looking observation and head off in one direction, but that a few frames later it becomes apparent that this is an error, and that the correct route goes another way. With an ordinary SISR particle filter the majority of the particles are now eliminated by resampling. The particle diversity of the state estimate will be very poor, and in a challenging problem there is a good chance that the track will be lost.

PICTURE

Improved estimates for previous state distributions can be made using the resample-move method of (Gilks & Berzuini 2001) to rejuvenate the particle diversity. However, this requires running a Markov chain for each particle on top of the usual importance sampling step, which can be computationally expensive.

ADD MORE COMPARISON WITH RESAMPLE-MOVE. IN SOME SENSE, THEY MUST BE EQUIVALENT - WORK OUT WHAT IT IS.

An alternative framework for fixed lag estimation is presented by (Doucet, Briers & Sénécal 2006). Now we propose new values not only for the current state but also the previous states within some time window. The effect of this in our tracking example is that we can “redirect” particles which have gone astray along the correct path, maintaining a better particle diversity and reducing the probability of track loss.

5.2 A mathematical framework for fixed lag estimation

Here we consider the mathematical framework for fixed-lag estimation, as presented by (Doucet et al. 2006) and (Briers, Doucet, Maskell & Horridge 2006).

As before, the target posterior distribution in which we are interested is the familiar $P(X_{1:t}|Y_{1:t})$. However, the proposal mechanism now becomes more complex, because we will be replacing states in an existing particle. We first propose a particle from which to take the state “history”, that is $X_{1:t-L}$. However, if we take this from the particle distribution from the previous step we get more of the path than we need, because each particle is a set of states $X_{1:t-1}$. The final $L-1$ states will be replaced when by a new “present”, $X'_{t-L+1:t}$, drawn from an importance distribution. The complete proposal is thus

$$\{X_{1:t-L}, X'_{t-L+1:t}\} \sim \int q(X_{1:t-1}|Y_{1:t-1})q(X'_{t-L+1}|X_{1:t-1}, Y_{t-L+1:t})dX_{t-L+1:t-1} \quad (5.1)$$

where $q(X_{1:t-1}|Y_{1:t-1})$ is a proposal distribution using the arbitrarily-weighted particles from $\hat{P}(X_{1:t-1}|Y_{1:t-1})$, in the auxiliary sampling sense. This integral is required for calculation of importance weights of acceptance probabilities, but in general it will be intractable. If we were to restrict our proposals to depend only on the history, i.e. use $q(X_{t-L+1:t}|X_{1:t-L}, Y_{t-L+1:t})$, then the proposal would become

$$\{X_{1:t-L}, X'_{t-L+1:t}\} \sim \hat{P}(X_{1:t-L}|Y_{1:t-1})q(X'_{t-L+1}|X_{1:t-1}, Y_{t-L+1:t}) \quad (5.2)$$

However, with this proposal, in order to evaluate importance weights or acceptance probabilities we still need to calculate

$$P(X_{1:t-L}|Y_{1:t-1}) = \frac{P(Y_{t-L+1:t-1}|X_{t-L}^{(m)})P(X_{1:t-L}|Y_{1:t-L})}{P(Y_{t-L+1:t-1}|Y_{1:t-L})} \quad (5.3)$$

The problematic term is:

$$P(Y_{t-L+1:t}|X_{t-L}^{(m)}) = \int P(Y_{t-L+1:t}|X_{t-L}^{(m)}, X_{t-L+1:t})P(X_{t-L+1:t}|X_{t-L}^{(m)})dX_{t-L+1:t} \quad (5.4)$$

This is intractable. The solution to this problem proposed in (Doucet et al. 2006) is to augment the dimension of the target distribution to include the discarded tracks. The new target is:

$$P(X_{1:t-L}, X'_{t-L+1:t}|Y_{1:t})\rho(X_{t-L+1:t-1}|X_{1:t-L}, X'_{t-L+1:t}) \quad (5.5)$$

Once a particle approximation has been generated for this target the required posterior distribution may be obtained by marginalisation. As we are simply going to discard the old track sections $X_{t-L+1:t-1}$, the choice of $\rho(\cdot)$ does not alter the distribution of the posterior. However, the variance of weights/acceptance probabilities may be affected.

With an expanded space for the target distribution there is no need to marginalise the previous states. The proposal distribution now becomes:

$$\{X_{1:t-1}, X'_{t-L+1:t}\} \sim q(X_{1:t-1}|Y_{1:t-1})q(X'_{t-L+1:t}|X_{1:t-1}, Y_{t-L+1:t}) \quad (5.6)$$

Histories are proposed from the $t - 1$ frame in the normal manner, see equation 3.28.

(Doucet et al. 2006) show that the optimum choice (in the minimal importance weight variance sense) of artificial conditional distribution is given by

$$\rho(X_{t-L+1:t-1}|X_{1:t-L}, X'_{t-L+1:t}) = P(X_{t-L+1:t-1}|X_{t-L}, Y_{t-L+1:t-1}) \quad (5.7)$$

with the current proposal distribution set to

$$q(X'_{t-L+1:t}|X_{1:t-s}, Y_{t-L+1:t}) = P(X'_{t-L+1:t}|X_{t-L}, Y_{t-L+1:t}) \quad (5.8)$$

Thus both artificial conditional and proposal distribution should equal to the conditional posterior. This is neither samplable nor calculable. However, it is intuitive to use the same approximation for each by making the artificial conditional equal to the proposal. The ratio of target to proposal probability is then given by:

$$\begin{aligned} & \frac{P(X_{1:t-L}, X'_{t-L+1:t}|Y_{1:t})\rho(X_{t-L+1:t-1}|X_{1:t-L}, X'_{t-L+1:t})}{q(X_{1:t-1}|Y_{1:t-1})q(X'_{t-L+1:t}|X_{1:t-1}, Y_{t-L+1:t})} \\ &= \frac{W_t}{V_t} \frac{P(X_{1:t-L}, X'_{t-L+1:t}|Y_{1:t})q(X_{t-L+1:t-1}|X_{1:t-L}, X'_{t-L+1:t-2}, Y_{t-L+1:t-1})}{P(X_{1:t-1}|Y_{1:t-1})q(X'_{t-L+1:t}|X_{1:t-1}, Y_{t-L+1:t})} \\ &\propto \frac{W_t}{V_t} \frac{P(Y_{t-L+1:t}|X'_{t-L+1:t})P(X'_{t-L+1:t}|X_{t-L})q(X_{t-L+1:t-1}|X_{1:t-L}, X'_{t-L+1:t-2}, Y_{t-L+1:t-1})}{P(Y_{t-L+1:t-1}|X_{t-L+1:t-1})P(X_{t-L+1:t-1}|X_{t-L})q(X'_{t-L+1:t}|X_{1:t-1}, Y_{t-L+1:t})} \end{aligned} \quad (5.9)$$

5.2.1 Proposing older histories

CONSIDER MOVING THIS TO THE MCMC IMPLEMENTATION SECTION So far we have chosen to propose histories from the $t - 1$ particle distribution. However, a suitable history could be also selected from the $t - s$ distribution, for $1 \leq s \leq L$. In a situation where a track has not been detected for several consecutive frames, taking a history from further in the past may make the proposal of a good candidate more probable.

PICTURE

The history proposal lag variable s can most easily be handled by sampling it too.

$$\{X_{1:t-s}, X'_{t-L+1:t}, s\} \sim q(s)q(X_{1:t-s}|Y_{1:t-s})q(X'_{t-L+1:t}|X_{1:t-s}, Y_{t-L+1:t}) \quad (5.10)$$

This is simply equivalent to enlarging the pool of particles from which we propose histories. $q(s)$ is simply a modification to the particle proposal weights.

SOMETHING ABOUT CONSERVATIVE RESAMPLING AND ITS EQUIVALENT AIM

5.3 Application to tracking

Erm... i've kinda said what I was going to say here in the previous section

5.3.1 Importance distributions

The dimensionality of a fixed lag particle filter will generally be very large. Consider a problem where each target has a 4-dimensional kinematic state (position and velocity in x and y) and an observation-association index. This gives us five dimensions per target per time step. If we use a lag window with length $L = 5$ and 5 targets we will have a 125 dimensional state. If we used a basic, bootstrap approach to such a problem, the chances of even a single particle following the correct path are negligible. Instead we must exploit the strong correlations between states arising from the structure of the problem. We first factorise the proposal thus:

$$\begin{aligned} q(x_{t-L+1:t}, \lambda_{t-L+1:t} | x_{t-L}, Y_{t-L+1:t}) \\ &= q(\lambda_{t-L+1:t} | x_{t-L}, Y_{t-L+1:t}) \\ &= q(x_{t-L+1:t} | x_{t-L}, \lambda_{t-L+1:t}, Y_{t-L+1:t}) \end{aligned} \quad (5.11)$$

Thus we can sample first the association variables then the state variables. Each of these terms can be further factorised over time, as we shall see. If we were to replace each of the factors with its corresponding posterior distribution then this factorisation would recreate the “optimal” importance distribution.

Association proposals

The proposal for the association variables is a discrete distribution over the possible observations with which the target could be associated. Within the factorisation above, the “optimal” form of the proposal is:

$$\begin{aligned} q(\lambda_{t-L+1:t} | x_{t-L}, Y_{t-L+1:t}) \\ &\propto P(Y_{t-L+1:t} | \lambda_{t-L+1:t}, x_{t-L}) P(\lambda_{t-L+1:t} | x_{t-L}) \\ &= \prod_{\substack{k=1:L \\ tt=t-L+k}} P(Y_{tt} | Y_{tt+1:t} \lambda_{tt:t}, x_{t-L}) P(\lambda_{tt}) \\ &= \prod_{\substack{k=1:L \\ tt=t-L+k}} \int P(Y_{tt} | x_{tt}, \lambda_k) P(x_{tt} | Y_{tt+1:t}, \lambda_{tt+1:t}, x_{t-L}) dx_{tt} P(\lambda_{tt}) \end{aligned} \quad (5.12)$$

This suggests a convenient sequential sampling procedure, starting with λ_t and working backwards in time. The normalisation constant is not known, but as this is discrete distribution we can enforce normalisation by dividing by the sum.

First we consider a factor from this expression with $\lambda_{tt} = 1$, i.e. a proposal that in a particular frame the target is not detected. In this case, the observation density is independent of the state of the target, and we have:

$$P(Y_{tt} | Y_{tt+1:t} \lambda_{tt:t}, x_{t-L}) P(\lambda_{tt}) = V^{-1} P(\lambda_t = 0) \quad (5.13)$$

When the target is detected, the factors of the proposal distribution of equation 5.12 may be calculated analytically for the linear-Gaussian case. For other cases, the EKF approximations may be used. As this is only a proposal distribution, such an approximation will not affect

the distribution of the particle distribution generated. The state distribution term over x_{tt} is problematic, requiring k integrals to calculate. Although this is will be analytic with Gaussian dynamics, its complexity may become unmanageable. This term represents the probability of the state given the set of observations associated with this target at later times. We can render this calculation more manageable by replacing the whole set of future observations with just one.

$$P(x_k|Y_{k+1:t}, \lambda_{k+1:t}, x_{t-L}) \approx P(x_k|Y_{k+d}, \lambda_{k+d}, x_{t-L}) \quad (5.14)$$

For cases with low observation noise, where an observation gives us significant information about the state of the target, this substitution will have little effect. Later observations cannot add much additional information. Again, as this is a proposal distribution, such a substitution will not affect the validity of the resulting particle distribution.

In general we will use $d = 1$, as the closest observation in time will give us the most information about x_{tt} . However, if the target is not detected at time $tt + 1$, then we can increase d to find the next detection of the target.

Using this approximation, for $\lambda_{t-L+k} \neq 0$, we have

$$P(Y_{tt}|Y_{tt+1:t}, \lambda_{tt:t}, x_{t-L})P(\lambda_{tt}) \propto \mathcal{N}(y_{tt}^{(\lambda_{tt})}|m_{tt}, S_{tt}) \quad (5.15)$$

where usually

$$S_{tt} = [I - R^{-1}C_{tt}\Sigma_{tt}C_{tt}R^{-1}]^{-1} \quad (5.16)$$

$$m_{tt} = SR^{-1}C_{tt}\Sigma_{tt}[(A^d)^TC_{tt+d}R_d^{-1}y_{tt+d}^{\lambda_{tt+d}} + Q_k^{-1}A^kx_{t-L}] \quad (5.17)$$

$$\Sigma_{tt} = [C_{tt}^TR^{-1}C_{tt} + (A^d)^TC_{tt+d}R_d^{-1}C_{tt+d}A^d + Q_k^{-1}]^{-1} \quad (5.18)$$

$$Q_d = \sum_{l=0}^{d-1} A^l Q(A^l)^T \quad (5.19)$$

$$R_d = R + C_{tt+d}Q_d(C_{tt+d})^T \quad (5.20)$$

We will need a different expression for the case when $tt = t$, because no future associations have yet been proposed. Similarly, if $tt < t$ but the future associations have all been proposed as missed detections, then there are no future observations to guide us, whatever choice of d we use. In these cases we have:

$$S_{tt} = R_d \quad (5.21)$$

$$m_{tt} = C_{tt}A^kx_{t-L} \quad (5.22)$$

For full derivations, see Appendix.

Finally, substituting for the association prior terms, we have:

$$q(\lambda_{t-L+1:t}|x_{t-L}, Y_{t-L+1:t}) \propto \prod_{\substack{k=1:L \\ tt=t-L+k}} \begin{cases} P_D \mathcal{N}(y_{tt}^{(\lambda_{tt})}|m_{tt}, S_{tt}) & \lambda_{tt} = 0 \\ (1 - P_D)\mu_C V^{-1} & \lambda_{tt} \neq 0 \end{cases} \quad (5.23)$$

This gives us a complete sequential mechanism for proposing the associations.

State proposals

Once the associations are fixed, the states can be proposed. When the state space model is linear-Gaussian, we can propose directly from the “optimal” importance distribution for the states using the forward-filtering-backward-sampling algorithm of (Chib 1996), as suggested in (Doucet et al. 2006). For nonlinear models, we can use EKF approximations, as for the associations. Once again, we factorise the proposal:

$$\begin{aligned}
 q(x_{t-L+1:t} | x_{t-L}, \lambda_{t-L+1:t}, Y_{t-L+1:t}) \\
 &= P(x_{t-L+1:t} | x_{t-L}, \lambda_{t-L+1:t}, Y_{t-L+1:t}) \\
 &= P(x_t | \lambda_{t-L+1:t}, Y_{t-L+1:t}, x_{t-L}) \prod_{k=t-L+1}^{t-1} P(x_k | \lambda_{t-L+1:k}, Y_{t-L+1:k}, x_{t-L}, x_{k+1}) \quad (5.24)
 \end{aligned}$$

where

$$P(x_k | \lambda_{t-L+1:k}, Y_{t-L+1:k}, x_{t-L}, x_{k+1}) \propto P(x_{k+1} | x_k) P(x_k | \lambda_{t-L+1:k}, Y_{t-L+1:k}, x_{t-L}) \quad (5.25)$$

The distributions $P(x_k | \lambda_{t-L+1:k}, Y_{t-L+1:k}, x_{t-L})$ are given by a Kalman filter, and are Gaussian with mean μ_k and covariance Σ_k . Thus the complete state proposal is given by:

$$q(x_{t-L+1:t} | x_{t-L}, \lambda_{t-L+1:t}, Y_{t-L+1:t}) = \mathcal{N}(x_t | \mu_t, \Sigma_t) \prod_{k=t-L+1}^{t-1} \mathcal{N}(x_k | m_k, S_k) \quad (5.26)$$

where

$$S_k = [A^T Q^{-1} A + \Sigma^{-1}]^{-1} \quad (5.27)$$

$$m_k = S_k [A^T Q^{-1} x_{k+1} + \Sigma^{-1} \mu_k] \quad (5.28)$$

Fixed Lag Particle Filters

6.1 Sequential Importance Sampling and Resampling

6.1.1 Coping with dimensionality

6.2 Markov Chain Monte Carlo

6.3 Marginalised Particle Filters

Plan Of Future Research

radar, insects, crowds, cells, jump-diffusions, cursor tracking...

Bibliography

- Al-Awadhi, F., Hurn, M. & Jennison, C. (2004). Improving the acceptance rate of reversible jump mcmc proposals, *Statistics & Probability Letters* **69**(2): 189 – 198.
URL: <http://www.sciencedirect.com/science/article/B6V1D-4CT4DWS-2/2/cab144a2f70e81f189a0e819c4517adc>
- Bar-Shalom, Y., Chang, K. & Blom, H. (1989). Automatic track formation in clutter with a recursive algorithm, *Decision and Control, 1989., Proceedings of the 28th IEEE Conference on*, IEEE, pp. 1402–1408.
- Bar-Shalom, Y., Daum, F. & Huang, J. (2009). The probabilistic data association filter, *Control Systems Magazine, IEEE* **29**(6): 82 –100.
- Bar-Shalom, Y. & Tse, E. (1975). Tracking in a cluttered environment with probabilistic data association, *Automatica* **11**(5): 451–460.
- Bayes, M. & Price, M. (1763). An essay towards solving a problem in the doctrine of chances. by the late rev. mr. bayes, frs communicated by mr. price, in a letter to john canton, amfrs, *Philosophical Transactions* **53**: 370.
- Blackman, S. (2004). Multiple hypothesis tracking for multiple target tracking, *IEEE Aerospace and Electronic Systems Magazine* **19**(1): 5–18.
- Blake, A., Isard, M. et al. (1998). *Active contours*, Vol. 8, Springer London.
- Briers, M., Doucet, A., Maskell, S. & Horridge, P. (2006). Fixed-lag sequential monte carlo data association, *Proceedings of SPIE*, Vol. 6236, p. 62360S. Paper copy from British Library.
- Bunch, P. & Godsill, S. (2010). Transcription of musical audio using poisson point processes and sequential mcmc, *Proc. of 7th International Symposium on Computer Music Modeling and Retrieval*.
- Cappé, O., Godsill, S. & Moulines, E. (2007). An overview of existing methods and recent advances in sequential monte carlo, *Proceedings of the IEEE* **95**(5): 899–924.

- Carmi, A., Godsill, S. & Septier, F. (2009). Evolutionary mcmc particle filtering for target cluster tracking, *IEEE 13th DSP Workshop and the 5th SPE Workshop*, Marco Island, Florida.
- Casella, G. & Robert, C. P. (1996). Rao-blackwellisation of sampling schemes, *Biometrika* **83**(1): 81–94.
URL: <http://biomet.oxfordjournals.org/content/83/1/81.abstract>
- Chib, S. (1996). Calculating posterior distributions and modal estimates in markov mixture models, *Journal of Econometrics* **75**(1): 79 – 97. Details forward-filtering-backward-sampling algorithm.
URL: <http://www.sciencedirect.com/science/article/B6VC0-3VWT1TW-7/2/b66dae744497c95ff3eddb744f301a6c>
- Daum, F. (2005). Nonlinear filters: beyond the kalman filter, *Aerospace and Electronic Systems Magazine, IEEE* **20**(8): 57 –69.
- Daum, F. & Huang, J. (2003). Curse of dimensionality and particle filters, *Aerospace Conference, 2003. Proceedings. 2003 IEEE*, Vol. 4, pp. 1979–1993.
- Doucet, A., Briers, M. & Sénécal, S. (2006). Efficient block sampling strategies for sequential monte carlo, *Journal of Computational and Graphical Statistics* **15**(3): 693–711.
- Doucet, A., Godsill, S. & Andrieu, C. (2000). On sequential monte carlo sampling methods for bayesian filtering, *Statistics and Computing* **10**: 197–208. 10.1023/A:1008935410038.
URL: <http://dx.doi.org/10.1023/A:1008935410038>
- Doucet, A. & Johansen, A. M. (2009). A tutorial on particle filtering and smoothing: Fifteen years later, in D. Crisan & B. Rozovsky (eds), *The Oxford Handbook of Nonlinear Filtering*, Oxford University Press.
URL: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.157.772&rep=rep1&type=pdf>
- Duane, S., Kennedy, A. D., Pendleton, B. J. & Roweth, D. (1987). Hybrid monte carlo, *Physics Letters B* **195**(2): 216 – 222.
URL: <http://www.sciencedirect.com/science/article/pii/037026938791197X>
- Fitzgerald, R. (1985). Track biases and coalescence with probabilistic data association, *Aerospace and Electronic Systems, IEEE Transactions on* **AES-21**(6): 822 –825.
- Fitzgerald, R. J. (1986). Development of practical pda logic for multitarget tracking by micro-processor, *American Control Conference, 1986*, pp. 889 –898.
- Fortmann, T., Bar-Shalom, Y. & Scheffe, M. (1983). Sonar tracking of multiple targets using joint probabilistic data association, *Oceanic Engineering, IEEE Journal of* **8**(3): 173 – 184.
- Geman, S. & Geman, D. (1984). Gibbs distributions, and the bayesian restoration of images, *IEEE Transactions on Pattern Analysis and Machine Intelligence* **6**(2): 721–741.

BIBLIOGRAPHY

- Gilks, W. R. & Berzuini, C. (2001). Following a moving target monte carlo inference for dynamic bayesian models, *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* **63**(1): 127–146.
URL: <http://dx.doi.org/10.1111/1467-9868.00280>
- Gilks, W., Richardson, S. & Spiegelhalter, D. (1996). *Markov chain Monte Carlo in practice*, Chapman and Hall.
- Girolami, M. & Calderhead, B. (2011). Riemann manifold langevin and hamiltonian monte carlo methods, *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* **73**(2): 123–214.
URL: <http://dx.doi.org/10.1111/j.1467-9868.2010.00765.x>
- Golightly, A. & Wilkinson, D. (2006). Bayesian sequential inference for nonlinear multivariate diffusions, *Statistics and Computing* **16**: 323–338. 10.1007/s11222-006-9392-x.
URL: <http://dx.doi.org/10.1007/s11222-006-9392-x>
- Gordon, N., Salmond, D. & Smith, A. (1993). Novel approach to nonlinear/non-gaussian bayesian state estimation, *Radar and Signal Processing, IEE Proceedings F* **140**(2): 107–113.
- Green, P. (1995). Reversible jump markov chain monte carlo computation and bayesian model determination, *Biometrika* **82**(4): 711.
- Hastings, W. K. (1970). Monte carlo sampling methods using markov chains and their applications, *Biometrika* **57**(1): 97–109.
URL: <http://biomet.oxfordjournals.org/content/57/1/97.abstract>
- Kalman, R. E. (1960). A new approach to linear filtering and prediction problems, *Journal Of Basic Engineering* **82**(Series D): 35–45.
URL: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.129.6247&rep=rep1&type=pdf>
- Khan, Z., Balch, T. & Dellaert, F. (2005). Mcmc-based particle filtering for tracking a variable number of interacting targets, *Pattern Analysis and Machine Intelligence, IEEE Transactions on* **27**(11): 1805–1819.
- Kitagawa, G. (1996). Monte carlo filter and smoother for non-gaussian nonlinear state space models, *Journal of Computational and Graphical Statistics* **5**(1): pp. 1–25.
URL: <http://www.jstor.org/stable/1390750>
- Laplace, P. S. (1774). Memoir on the probability of the causes of events, *Statistical Science* **1**(3): pp. 364–378.
URL: <http://www.jstor.org/stable/2245476>
- Liu, J. S. & Chen, R. (1995). Blind deconvolution via sequential imputations, *Journal of the American Statistical Association* **90**(430): pp. 567–576.
URL: <http://www.jstor.org/stable/2291068>
- Mahler, R. (1994). Random-set approach to data fusion, *Proceedings of SPIE*, Vol. 2234, p. 287.

- Mahler, R. (2003). Multitarget bayes filtering via first-order multitarget moments, *Aerospace and Electronic Systems, IEEE Transactions on* **39**(4): 1152 – 1178.
- Mahler, R. (2004). "statistics 101" for multisensor, multitarget data fusion, *Aerospace and Electronic Systems Magazine, IEEE* **19**(1): 53 –64.
- Metropolis, N., Rosenbluth, A., Rosenbluth, M., Teller, A., Teller, E. et al. (1953). Equation of state calculations by fast computing machines, *The journal of chemical physics* **21**(6): 1087.
- Musicki, D. & Evans, R. (2004). Joint integrated probabilistic data association: Jipda, *Aerospace and Electronic Systems, IEEE Transactions on* **40**(3): 1093 – 1099.
- Musicki, D., Evans, R. & Stankovic, S. (1994). Integrated probabilistic data association, *Automatic Control, IEEE Transactions on* **39**(6): 1237 –1241.
- Oh, S., Russell, S. & Sastry, S. (2004). Markov chain monte carlo data association for general multiple-target tracking problems, *Decision and Control, 2004. CDC. 43rd IEEE Conference on*, Vol. 1, pp. 735 –742 Vol.1.
- Oh, S., Russell, S. & Sastry, S. (2009). Markov chain monte carlo data association for multi-target tracking, *Automatic Control, IEEE Transactions on* **54**(3): 481 –497.
- Pang, S., Godsill, S., Li, J. & Septier, F. (2011). Sequential inference for dynamically evolving groups of objects, in D. Barber, A. Cemgil & S. Chiappa (eds), *Inference and Learning in Dynamic Models*, CUP. To Appear.
- Pang, S. K., Li, J. & Godsill, S. (2008). Models and algorithms for detection and tracking of coordinated groups, *Aerospace Conference, 2008 IEEE*, pp. 1 –17.
- Pitt, M. K. & Shephard, N. (1999). Filtering via simulation: Auxiliary particle filters, *Journal of the American Statistical Association* **94**(446): pp. 590–599.
URL: <http://www.jstor.org/stable/2670179>
- Rauch, H., Tung, F. & Striebel, C. (1965). Maximum likelihood estimates of linear dynamic systems, *AIAA journal* **3**(8): 1445–1450.
- Reid, D. (1979). An algorithm for tracking multiple targets, *Automatic Control, IEEE Transactions on* **24**(6): 843 – 854.
- Salmond, D. (1990). Mixture reduction algorithms for target tracking in clutter, *Proceedings of SPIE*, Vol. 1305, p. 434.
- Särkkä, S., Vehtari, A. & Lampinen, J. (2007). Rao-blackwellized particle filter for multiple target tracking, *Information Fusion* **8**(1): 2 – 15. Special Issue on the Seventh International Conference on Information Fusion-Part II, Seventh International Conference on Information Fusion.
URL: <http://www.sciencedirect.com/science/article/B6W76-4HJRRX5-1/2/c4d5356937753f21dc812b7943ffd6cd>
- Sea, R. G. (1971). An efficient suboptimal decision procedure for associating sensor data with stored tracks in real-time surveillance systems, *Decision and Control, 1971 IEEE Conference on*, Vol. 10, pp. 33 –37.

- Septier, F., Pang, S. K., Carmi, A. & Godsill, S. (2009). On mcmc-based particle methods for bayesian filtering: Application to multitarget tracking, *Computational Advances in Multi-Sensor Adaptive Processing (CAMSAP), 2009 3rd IEEE International Workshop on*, pp. 360 –363.
- Singer, R., Sea, R. & Housewright, K. (1974). Derivation and evaluation of improved tracking filter for use in dense multitarget environments, *Information Theory, IEEE Transactions on* **20**(4): 423 – 432.
- Streit, R. & Luginbuhl, T. (1994). Maximum likelihood method for probabilistic multihypothesis tracking, *Proceedings of SPIE*, Vol. 2235, p. 394.
- Svensson, L., Svensson, D. & Willett, P. (2009). Set jpda algorithm for tracking unordered sets of targets, *Information Fusion, 2009. FUSION '09. 12th International Conference on*, pp. 1187 –1194.
- Vermaak, J., Godsill, S. & Perez, P. (2005). Monte carlo filtering for multi target tracking and data association, *Aerospace and Electronic Systems, IEEE Transactions on* **41**(1): 309 – 332.
- Vo, B.-N. & Ma, W.-K. (2006). The gaussian mixture probability hypothesis density filter, *Signal Processing, IEEE Transactions on* **54**(11): 4091 –4104.
- Vo, B.-N., Singh, S. & Doucet, A. (2005). Sequential monte carlo methods for multitarget filtering with random finite sets, *Aerospace and Electronic Systems, IEEE Transactions on* **41**(4): 1224 – 1245.
- Whiteley, N., Singh, S. & Godsill, S. (2010). Auxiliary particle implementation of probability hypothesis density filter, *Aerospace and Electronic Systems, IEEE Transactions on* **46**(3): 1437 –1454.
- Willett, P., Ruan, Y. & Streit, R. (2002). Pmht: problems and some solutions, *Aerospace and Electronic Systems, IEEE Transactions on* **38**(3): 738 – 754.
- Wood, T. (2010). Random finite set theory for tracking.
URL: http://people.maths.ox.ac.uk/~woodtm/rfs_overview.pdf